

MATH 292.01

November 13, 2025

Taehun Cha
(Teaching Fellow)

Contents

- Recap) Gradient Descent
- Vanishing Gradient
- Functional Data Analysis
- Curse of Smoothness
 - We will observe additional difficulty in NN training



Recap) Gradient Descent

- For a function f with input variable x ,
 - Gradient Descent: $x' = x - \gamma \nabla_x f$
- For a loss function L with data $\mathcal{D} = \{(x, y)\}_{i=1}^N$ and weight w
 - Gradient Descent: $w' = w - \gamma \nabla_w L(x, y)$
 - For multiple data points, simply: $w' = w - \gamma \frac{1}{N} \sum_i \nabla_w L(x_i, y_i)$
- Sampling the small portion of data works nicely and efficiently
 - Stochastic Gradient Descent: $w' = w - \gamma \frac{1}{|B|} \sum_{i \in B} \nabla_w L(x_i, y_i)$, B: minibatch

Recap) Gradient Descent

- Theoretically, under the following conditions:
 - f is convex, (no saddle, no local minima)
 - f is L -smooth, (it does not ‘spike’ suddenly)
 - and learning rate γ is small enough, ($< \frac{1}{L}$)
- GD results in $x^* = \arg_x \min f(x)$ with infinite steps
- Does this hold in neural network training?



Recap) Gradient Descent

- In neural network training:
 - f is NOT convex, (bunch of saddle and local minima)
 - we don't know whether f is L -smooth, (Loss can ‘spikes’ suddenly)
 - (at least) we can set γ is small enough, ($< \frac{1}{L}$)
 - But we don't know L !!
- NN loss landscape is too wild 😱
- But surprisingly, it works nicely in many cases! 🤔

Recap) Gradient Descent

모집일정
2026학년도 전기 대학원 입학전형 일정 및 지원자격



전형별일정

구분	2026학년도 전기 원서접수 일정
내국인 전형	신입학 전형
	인터넷 원서접수 2025년 09월 26(금) - 10월 02일(목)
외국인 전형	신입학 전형
	인터넷 원서접수 2025년 09월 01(월) - 09월 10일(수)
학·석사연계과정 전형	2025년 12월 1일 - 12월 3일
석·박사통합과정 (진입)전형	2025년 12월 1일 - 12월 3일

Theories exist but under strict conditions...

In other word... many things to discover yet!



Curse of Smoothness in Functional Neural Networks

Taehun Cha and Donghun Lee

Abstract—Functional neural networks (FNNs) have emerged as powerful tools for modeling complex relationships in functional data, leveraging the flexibility of deep learning to capture non-linear patterns. However, most components of FNNs are directly

traditional data analysis—such as dimension reduction, regression, and classification—all specifically adapted to preserve the infinite-dimensional nature of functional observations.



AIML@K

Vanishing Gradient

- Assume an L -layer neural network f with
 - Hidden states at l -th layer: $h_l = \sigma(W_{l-1}h_{l-1})$
 - σ : non-linear activation function
 - W_{l-1} : weight matrix of $l - 1$ -th layer
 - $h_0 = x$, i.e. input data
- We train the NN with a loss function \mathcal{L}



Vanishing Gradient

- What we know
 - Train with GD, $W_l \leftarrow W_l - \gamma \nabla_{W_l} \mathcal{L}$
 - Chain rule: $\nabla_{W_l} \mathcal{L} = (\nabla_{h_L} \mathcal{L}) \cdot (\nabla_{h_{L-1}} h_L) \cdot (\dots) \cdot (\nabla_{h_{l+1}} h_{l+2}) \cdot (\nabla_{W_l} h_{l+1})$
 - where $\nabla_{h_{l+1}} h_{l+2} = \text{diag}[\sigma'(W_{l+1}h_{l+1})] \cdot W_{l+1}$, σ' : derivative of σ
 - and $\nabla_{[W_l]_{i,:}} [h_{l+1}]_i = [\sigma'(W_l h_l)]_i \cdot h_l$,
 - (Big Beautiful) Backpropagation
 - Compute the loss value $\mathcal{L}(x, y; f)$
 - Compute the gradient backward, $\nabla_{h_L} \mathcal{L}, \nabla_{h_{L-1}} h_L, \nabla_{h_{l+1}} h_{l+2}, \dots$
 - Compute $\nabla_{W_l} h_{l+1}$
 - Then multiply!



AIML@K

Vanishing Gradient

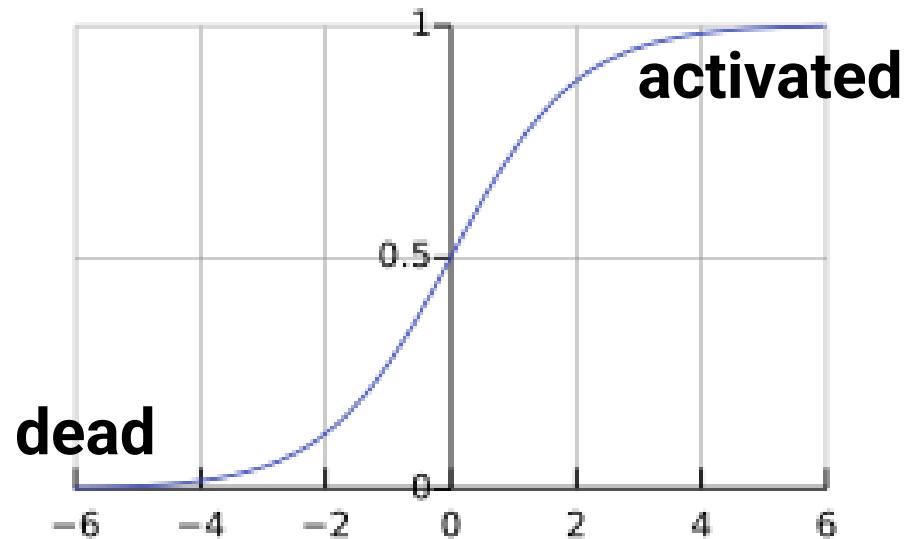
- Recently, deeeeep NN shows much better performance
 - GPT-4 is estimated to use 120 layers
 - ResNet-152 uses 152 layers
- But the theoretical benefit of depth is not yet established
- Moreover, deeeeep NN suffers from the **vanishing gradient**



Vanishing Gradient

- Suppose σ is Sigmoid function, i.e.

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

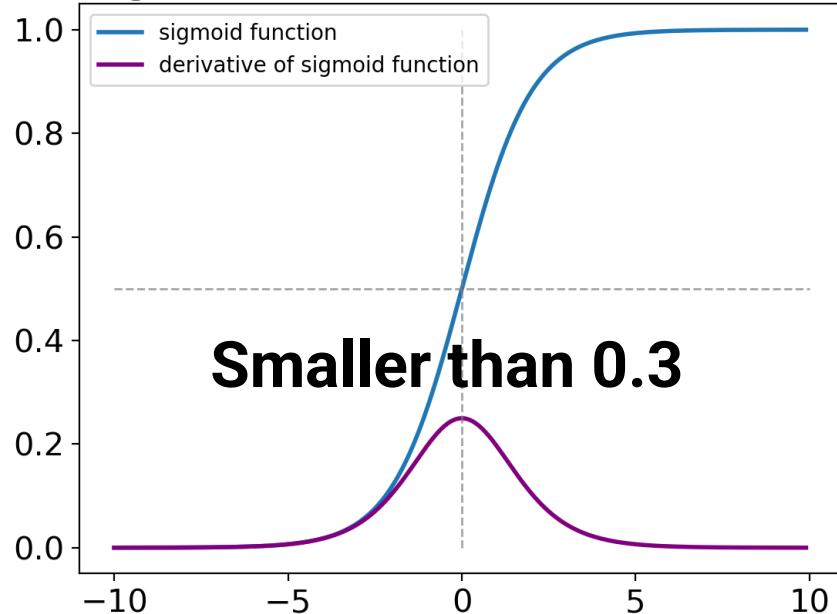


Vanishing Gradient

- What if we draw σ' ? i.e.

$$\sigma'(x) = \sigma(x)[1 - \sigma(x)]$$

sigmoid function and its derivative



and (near) zero for too many part of domain

Vanishing Gradient

- What if we multiply the multiple $\sigma' \in (0, 0.3)$?

- $0.3^{10} = 0.000006$

- Recall the chain rule,

$$\nabla_{W_l} \mathcal{L} = (\nabla_{h_L} \mathcal{L}) \cdot (\nabla_{h_{L-1}} h_L) \cdot (\dots) \cdot (\nabla_{h_{l+1}} h_{l+2}) \cdot (\nabla_{W_l} h_{l+1})$$

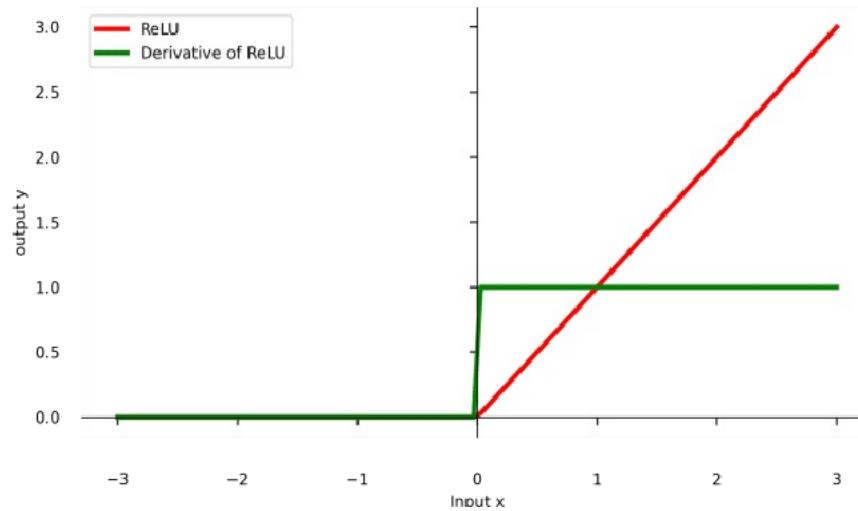
where $\nabla_{h_{l+1}} h_{l+2} = \text{diag}[\sigma'(W_{l+1} h_{l+1})] \cdot W_{l+1}$

σ' everywhere!

- As a result, W_l with low l would receive nearly 0 gradient

Vanishing Gradient

- How the modern NN resolved this problem?
 - Activation functions (e.g. ReLU) with higher derivative values



Vanishing Gradient

- How the modern NN resolved this problem?
 - Activation functions (e.g. ReLU) with higher derivative values
 - Residual connection:

$$h_l = \sigma(W_{l-1}h_{l-1}) + h_{l-1}$$

$$\nabla_{h_{l-1}} h_l = \text{diag}[\sigma'(W_{l-1}h_{l-1})] \cdot W_{l-1} + I$$

- ResNet is Res(idual)Net
- GPT (Transformers) also use this



AIML@K

Functional Data Analysis

- There is a “function”
 - Think it as an infinite-dimensional vector
 - Think of a vector $a = [1, 2, 3]$, where $a_i = i$.
 - Now stretch it two 5 dimension $[1, 1.5, 2, 2.5, 3]$.
 - Stretch it into 9 dimension $[1, 1.25, 1.5, 1.75, 2, 2.25, 2.5, 2.75, 3]$.
 - ...
 - Then we can obtain a vector looks like $f(x) = x, x \in [1, 3]$



Functional Data Analysis

- There is a “function”
- We are familiar with finite-dimension data
- But there exists infinite-dimension one
 - Continuous timeseries, electric signal, voice, ...
- Classic statistics preprocessed it into a low-dimensional data
- Some statisticians wanted to **preserve** the functional property

Functional Data Analysis

- Basic example: Linear regression
 - What we know: $y = w^\top x, x \in \mathbf{R}^d$
 - What FDA do: $y = \int \beta(t)x(t)dt$
- Why we need this?
 - If we represent the weight function $\beta(t) = \sum_i w_i \phi_i(t)$,
 - where ϕ : basis function (Fourier, Legendre, ...)
 - If we use 5 basis, we do LR of inf-dim function with only 5 parameters!



Functional Data Analysis

- LR to Neural Network
 - Recall the multi-target LR: $y = Wx$,
 - Recall the NN: $h_l = \sigma(W_{l-1}h_{l-1})$
 - It can be seen as a recursive LR + non-linear σ



Functional Data Analysis

- Functional LR to Functional Neural Network

- Recall the functional LR: $y = \int \beta(t)x(t)dt,$
- (Univariate) FNN: $h_l^k(t) = \sigma\left(\sum_j \beta_{l-1}^{j,k}(t) h_{l-1}^j(t)\right),$
- (Bivariate) FNN: $h_l^k(t) = \sigma\left(\sum_j \int \beta_{l-1}^{j,k}(s, t) h_{l-1}^j(s) ds\right),$
 - j : previous dimension, k : current dimension
 - and bivariate weight function $\beta(s, t) = \sum_i w_i \phi_i(s) w_j \psi_j(t)$
- It can be seen as a recursive FLR + non-linear σ



AIML@K

Curse of Smoothness

- FDA is proposed to overcome the *Curse of Dimensionality*
 - Input data is infinite-dimension → need infinite weight?
 - Not with FDA!
- But there was another curse... *Curse of Smoothness...* 😱

Curse of Smoothness

- Initial observation
 - I found FNN works so poorly when it gets deeper
 - When σ : Sigmoid, deeper DNN and FNN perform worse
 - I first thought it's just vanishing gradient

Dataset	Model	Sigmoid			
		2 layer	4 layer	6 layer	Δ Acc.
EEG	DNN	55.44% (1.74)	53.63% (0.96)	46.72% (0.88)	-8.72
	U-FNN	69.98% (2.11)	70.48% (2.26)	65.97% (2.23)	-4.01
	B-FNN	69.68% (1.06)	66.09% (1.77)	29.82% (6.12)	-39.86

Curse of Smoothness

- Initial observation
 - I found FNN works so poorly when it gets deeper
 - When σ : Sigmoid, deeper DNN and FNN perform worse
 - I first thought it's just vanishing gradient
 - But it persists with ReLU!
 - Though DNN recovered!

Dataset	Model	Sigmoid				ReLU			
		2 layer	4 layer	6 layer	Δ Acc.	2 layer	4 layer	6 layer	Δ Acc.
EEG	DNN	55.44% (1.74)	53.63% (0.96)	46.72% (0.88)	-8.72	52.24% (1.18)	53.70% (1.37)	53.62% (1.56)	+1.38
	U-FNN	69.98% (2.11)	70.48% (2.26)	65.97% (2.23)	-4.01	69.99% (1.52)	65.51% (1.33)	65.58% (2.02)	-4.41
	B-FNN	69.68% (1.06)	66.09% (1.77)	29.82% (6.12)	-39.86	68.21% (1.49)	63.93% (2.35)	52.35% (4.70)	-15.86

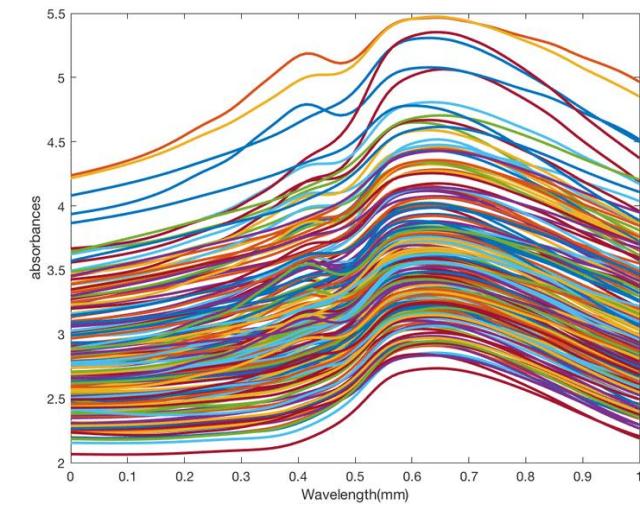
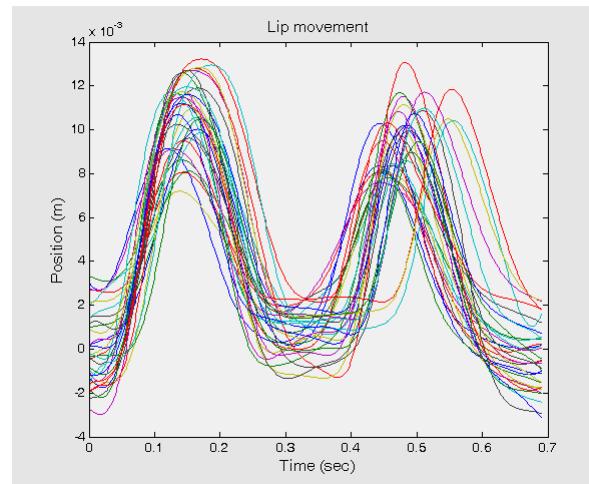
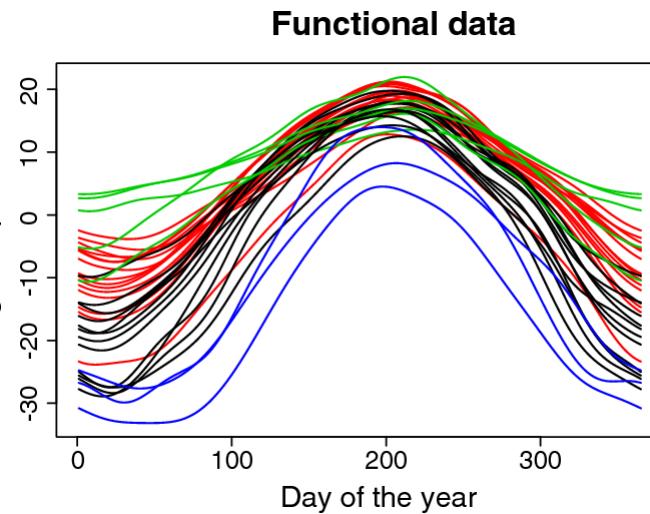
Curse of Smoothness

- Initial observation
 - It happens everywhere!
 - Why does it happen?
 - Even with ReLU?

Dataset	Model	Sigmoid				ReLU			
		2 layer	4 layer	6 layer	Δ Acc.	2 layer	4 layer	6 layer	Δ Acc.
Growth	DNN	95.67% (1.63)	94.62% (1.63)	96.72% (1.33)	+1.05	96.72% (0.0)	95.67% (0.0)	94.62% (0.0)	-2.10
	U-FNN	96.72% (1.63)	81.87% (4.52)	51.52% (6.18)	-45.20	96.72% (1.63)	93.56% (5.33)	90.41% (2.50)	-6.31
	B-FNN	95.67% (1.33)	76.55% (7.48)	80.70% (6.18)	-14.97	95.67% (1.33)	86.20% (2.50)	79.71% (4.52)	-15.96
Tecator	DNN	93.95% (0.71)	94.88% (1.69)	93.95% (1.30)	0.0	93.95% (1.59)	91.16% (1.92)	92.56% (2.49)	-1.39
	U-FNN	88.37% (3.71)	66.05% (4.14)	64.19% (4.43)	-24.18	93.95% (0.92)	86.51% (6.61)	78.14% (2.98)	-15.81
	B-FNN	87.44% (2.52)	68.84% (6.28)	65.58% (5.14)	-21.86	93.02% (1.42)	83.72% (5.83)	60.93% (4.67)	-32.09
EEG	DNN	55.44% (1.74)	53.63% (0.96)	46.72% (0.88)	-8.72	52.24% (1.18)	53.70% (1.37)	53.62% (1.56)	+1.38
	U-FNN	69.98% (2.11)	70.48% (2.26)	65.97% (2.23)	-4.01	69.99% (1.52)	65.51% (1.33)	65.58% (2.02)	-4.41
	B-FNN	69.68% (1.06)	66.09% (1.77)	29.82% (6.12)	-39.86	68.21% (1.49)	63.93% (2.35)	52.35% (4.70)	-15.86

Curse of Smoothness

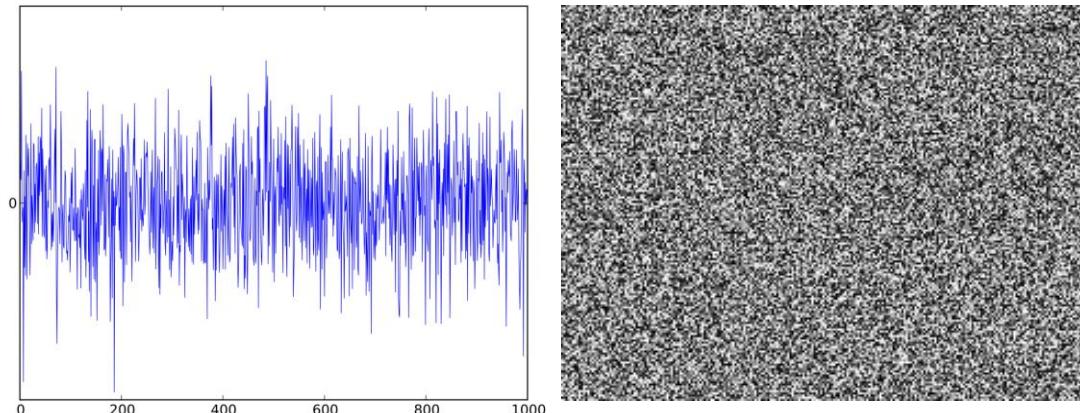
- Hypothesis: may be the functional smoothness is cause
 - FDA is proposed to represent the **smooth** functions



Functional Data Examples

Curse of Smoothness

- Hypothesis: may be the functional smoothness is cause
 - FDA is proposed to represent the **smooth** functions
 - Usually don't say white noise is functional data



Non-Functional Data Examples



Curse of Smoothness

- Recall, the definition of L -smoothness
 - “ f is L -smooth” is equiv. to “ ∇f is L -Lipschitz”
 - Function $f: \mathcal{X} \rightarrow \mathcal{Y}$ is L -Lipschitz if $\forall x', x: \|f(x') - f(x)\| \leq L\|x' - x\|$
- Where we used it?

What Happens in 1-step?

- 1-step: $x' \leftarrow x - \gamma \nabla f(x)$
- Start with Taylor expansion of f at x'
- $f(x') \approx f(x) + \nabla f(x)^\top (x' - x) + \frac{1}{2}(x' - x)^\top \nabla^2 f(x)(x' - x)$
 - Claim1: $x' - x = -\gamma \nabla f(x)$
 - True, by “1-step” of GD
 - Claim2: $\frac{1}{2}(x' - x)^\top \nabla^2 f(x)(x' - x) \leq \frac{1}{2}L\|x' - x\|_2^2$
 - True, by f is L -smooth \Rightarrow bounded eigenvalues of Hessian of f (proof?)
 - Actually, this approach requires f to be twice differentiable (to have Hessian)



10/28/2025

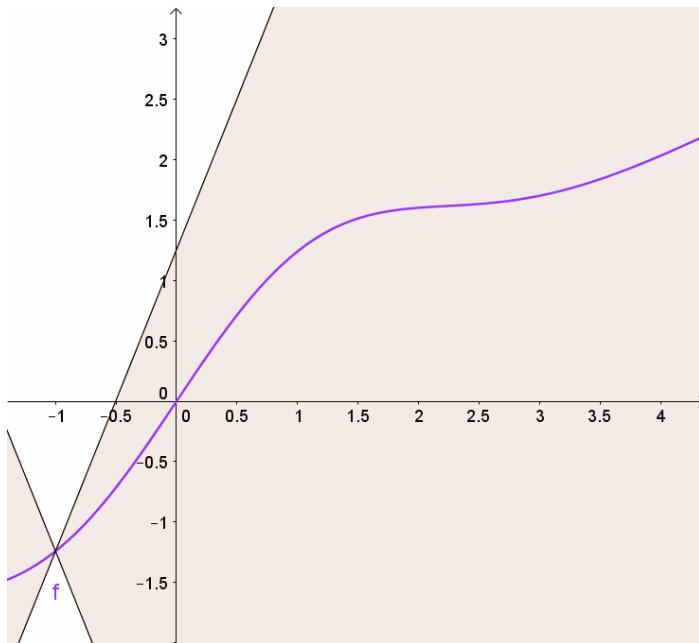
MATH 292.01

8



Curse of Smoothness

- We will use the L -Lipschitz on f , not ∇f
 - It's not " L -Smooth", but it also implies some smoothness



Curse of Smoothness

- Settings
 - Let a L-layer functional NN with hidden $h_l(t) = \sigma(\beta_{l-1}(t)h_{l-1}(t))$,
 - where $\beta_{l-1}(t) = \sum_b w_b \phi_b(t)$.
 - It's univariate, single hidden dimension (but can be extended).
- Assume,
 - $h_{l-1}(t)$ and $\phi_i(t)$ are l_1 and l_2 -Lipschitz (Smooth),
 - $h_{l-1}(t)$ and $\phi_i(t)$ are zero at some point (Normalization),
 - All functions are in $L^1(0,1)$, $L^2(0,1)$, and $L^4(0,1)$ (Another smooth).



AIML@K

Curse of Smoothness

- Assume,
 - $h_{l-1}(t)$ and $\phi_i(t)$ are l_1 and l_2 -Lipschitz (Smooth),
 - $h_{l-1}(t)$ and $\phi_i(t)$ are zero at some point (Normalization),
 - All functions are in $L^1(0,1)$, $L^2(0,1)$, and $L^4(0,1)$ (Another smooth).
- Statement:
 - Then $\left\| \frac{dh_l}{dw_b} \right\|_1 \leq \sqrt{0.2} \cdot \left\| |\sigma'(\beta_{l-1}(t)h_{l-1}(t))| \right\|_2 \cdot l_1 \cdot l_2$



Curse of Smoothness

- Recall the Vanishing Gradient

Vanishing Gradient

- What if we multiply the multiple $\sigma' \in (0, 0.3)$?
 - $0.3^{10} = 0.000006$
- Recall the chain rule,
$$\nabla_{W_l} \mathcal{L} = (\nabla_{h_L} \mathcal{L}) \cdot (\nabla_{h_{L-1}} h_L) \cdot (\dots) \cdot (\nabla_{h_{l+1}} h_{l+2}) \cdot (\nabla_{W_l} h_{l+1})$$

where $\nabla_{h_{l+1}} h_{l+2} = \text{diag}[\sigma'(W_{l+1} h_{l+1})] \cdot W_{l+1}$
 σ' everywhere!
- As a result, W_l with low l would receive nearly 0 gradient

11/13/2025 MATH 292.01 (TF Session) 13 KOREA UNIVERSITY

- Now check the statement again:

If $\sqrt{0.2} \cdot l_1 \cdot l_2 < 1 \dots$

• Then $\left\| \frac{dh_l}{dw_b} \right\|_1 \leq \sqrt{0.2} \cdot \left\| \sigma'(\beta_{l-1}(t)h_{l-1}(t)) \right\|_2 \cdot l_1 \cdot l_2$

Multiplied L times

Resolved with ReLU

Curse of Smoothness

- So, the statement,
 - $\left\| \frac{dh_l}{dw_b} \right\|_1 \leq \sqrt{0.2} \cdot \left\| \sigma'(\beta_{l-1}(t)h_{l-1}(t)) \right\|_2 \cdot l_1 \cdot l_2$
 - implies gradient norm is bounded by the function smoothness
 - Lower Lipschitz Constant => Smoother
 - And it make gradient vanish at deeper layer
 - **Curse of Smoothness!!**

Curse of Smoothness

- Cf) Short Proof

$$\begin{aligned} \bullet \quad & \left\| \frac{dh_l}{dw_b} \right\| = \left\| \sigma'(\beta_{l-1}(t)h_{l-1}(t)) \cdot h_{l-1}(t) \cdot \phi_b(t) \right\| \\ & \leq \left\| \sigma'(\beta_{l-1}(t)h_{l-1}(t)) \right\|_2 \cdot \left\| h_{l-1}(t) \right\|_4 \cdot \left\| \phi_b(t) \right\|_4 \text{ by generalized Holder} \end{aligned}$$

- Note, if a function $g \in L^4(0, 1)$ is L -Lipschitz and zero at t' , then

$$|g(t')| = |g(t') - g(t)| \leq L|t' - t|$$

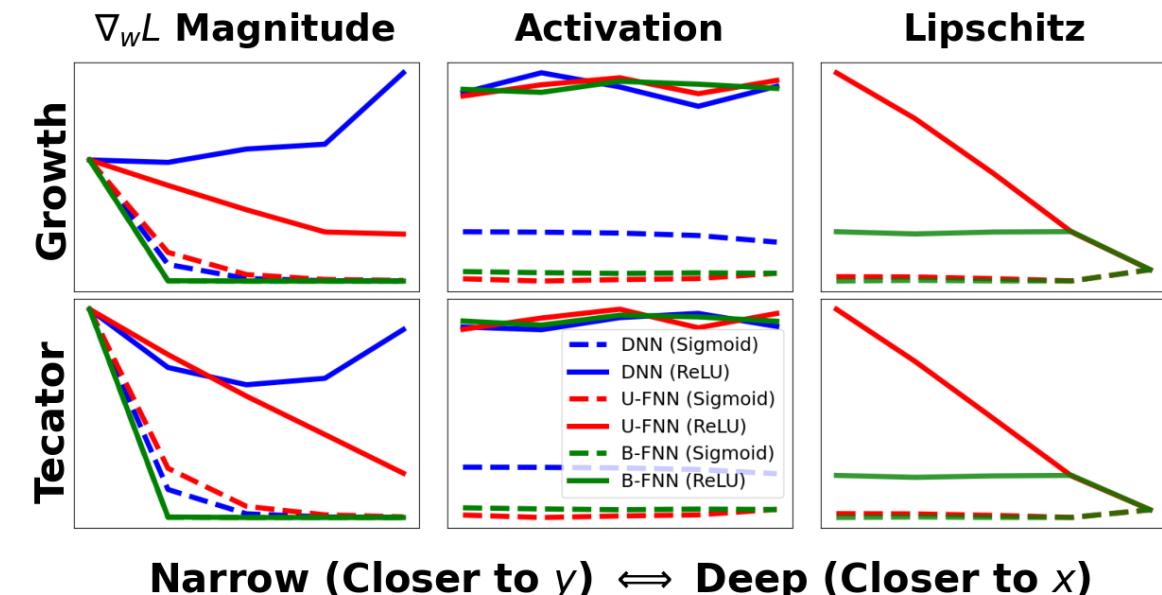
- and $\|g(t)\|_4 = \left[\int_{(0,1)} |g(t)|^4 dt \right]^{1/4} \leq L \left[\int_{(0,1)} |t' - t|^4 dt \right]^{1/4} \leq L \left[\frac{1}{5} \right]^{1/4}$ (check!)
- So $\|h_{l-1}(t)\|_4 \leq l_1 \left[\frac{1}{5} \right]^{1/4}$ and $\|\phi_b(t)\|_4 \leq l_2 \left[\frac{1}{5} \right]^{1/4}$



AIML@K

Curse of Smoothness

- Empirical Results
 - With Sigmoid (dashed line)
 - All models suffer from vanishing grad
 - with low activation gradient σ'
 - + In FNN, smoothness gets severe
 - With ReLU (real line)
 - DNN recovered but FNN couldn't
 - Despite high activation gradient σ' ,
 - Smoothness hinder it to obtain high grad.



Conclusion

- Optimizing NN with GD is so difficult
 - Suffer from non-convex, non-smooth loss landscape
 - Suffer from **vanishing gradient** (maybe solved)
- Statisticians wanted to overcome the **Curse of Dimensionality**
 - in infinite-dimensional functional data
 - with functional data analysis framework
- However, functional data falls into the **Curse of Smoothness**
 - Which is functional version of **vanishing gradient**

Conclusion

- Optimizing NN with GD is so difficult
 - Suffer from non-convex, non-smooth loss landscape
 - Suffer from **vanishing gradient** (maybe solved)] **We thought it's resolved**
- Statisticians wanted to overcome the **Curse of Dimensionality**] **Different problem appeared**
 - in infinite-dimensional functional data
 - with functional data analysis framework] **Solution proposed**
- However, functional data falls into the **Curse of Smoothness**] **Another Problem showed up**
 - Which is functional version of **vanishing gradient**] **Related to previously resolved issue**



AIML@K